



---

# 4CS372: Advance Database System Lab (ADSL)

---

Assignment NO:9



May 21, 2021

GAURAV ARORA(PRN:2018BTECS00032)

VIDHAN SHAH(PRN:2018BTECS00058) ,

**GroupNo:2018BCGRP24**

# Title: Cassandra Clustering

## Objective/Aim:

1. Setup a multi-node Cassandra Cluster on single windows machine.
2. Install the DataStax OpsCenter community edition.
3. Demonstrate the cluster operations of above use case using OpsCenter.

## Introduction & Theory:

Apache Cassandra is an open source NoSQL distributed database trusted by thousands of companies for scalability and high availability without compromising performance. Linear scalability and proven fault-tolerance on commodity hardware or cloud infrastructure make it the perfect platform for mission-critical data.

What is Cluster?

The cluster is a collection of nodes that represents a single system. A cluster in Cassandra is one of the shells in the whole Cassandra database. Many Cassandra Clusters combine together to form the database in Cassandra. A Cluster is basically the outermost shell or storage unit in a database. The Cassandra Cluster contains many different layers of storage units. Each layer contains the other.

## Implementation:

**1) Setup a multi-node Cassandra Cluster on single windows machine. Give your group name (2018BCGRP\*\*) to cluster. Follow the steps given in below link <https://extendit.us/articles/steps-configure-multiple-nodescassandra-single-windows-machine>**

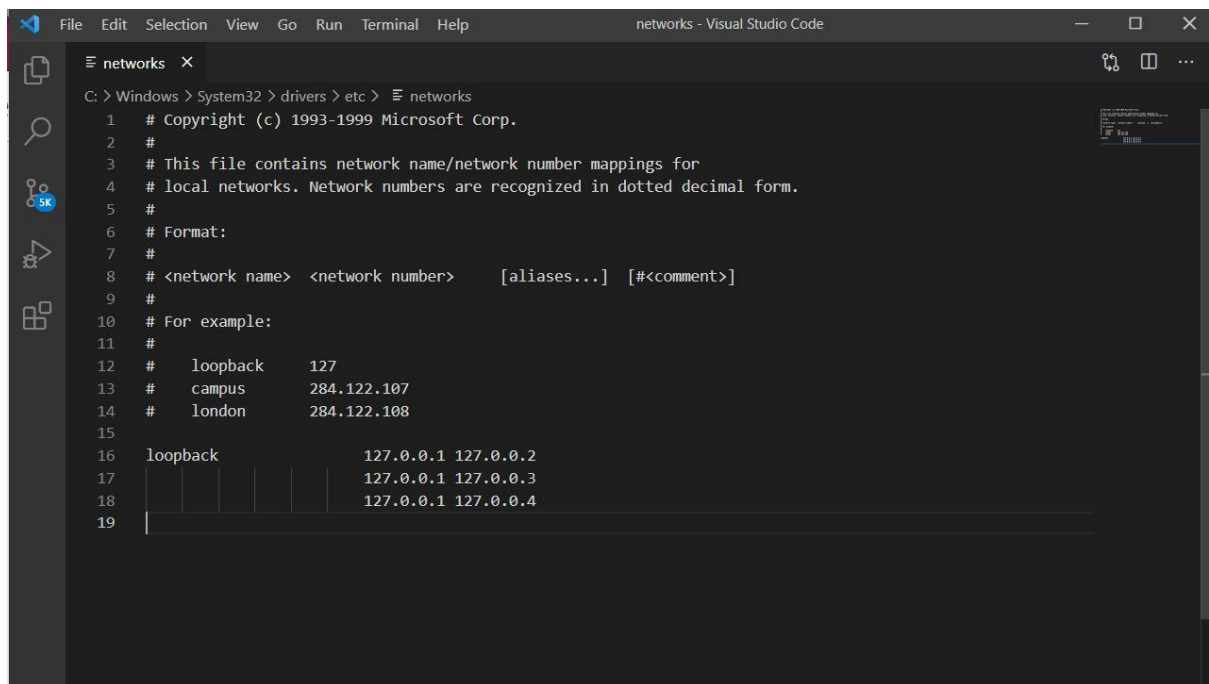
Step 1:

C:\Windows\System32\drivers\etc

Update the file with the following entries:

127.0.0.1 127.0.0.2  
127.0.0.1 127.0.0.3  
127.0.0.1 127.0.0.4

## 2018BCGRP24



```
C: > Windows > System32 > drivers > etc > networks

1 # Copyright (c) 1993-1999 Microsoft Corp.
2 #
3 # This file contains network name/network number mappings for
4 # local networks. Network numbers are recognized in dotted decimal form.
5 #
6 # Format:
7 #
8 # <network name> <network number> [aliases...] [#<comment>]
9 #
10 # For example:
11 #
12 #     loopback      127
13 #     campus        284.122.107
14 #     london        284.122.108
15 #
16 loopback          127.0.0.1 127.0.0.2
17                  127.0.0.1 127.0.0.3
18                  127.0.0.1 127.0.0.4
19
```

### Step 2:

Next, create the folder structure for your Cassandra installation.

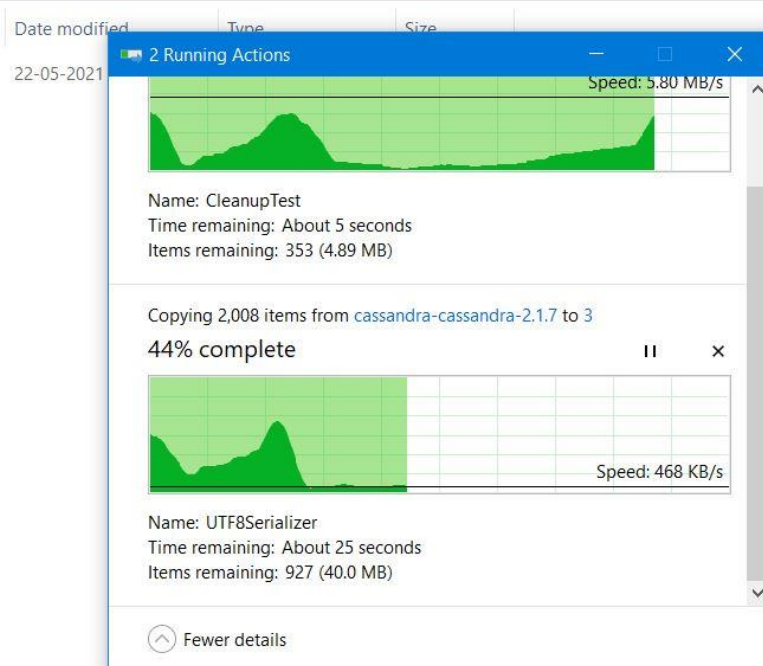
Installation: Install the DataStax OpsCenter community edition



The screenshot shows the Apache Cassandra download page. At the top, there is a navigation bar with links: News, About, Make a Donation, The Apache Way, Join Us, and Downloads. Below this is the Apache Software Foundation logo and the text "COMMUNITY-LED DEVELOPMENT 'THE APACHE WAY'". A secondary navigation bar includes links for Projects, People, Community, License, and Sponsors. The main content area states: "We suggest the following mirror site for your download: <https://downloads.apache.org/cassandra/3.11.10/apache-cassandra-3.11.10-bin.tar.gz>". It also mentions other mirror sites and provides instructions on how to verify the integrity of the downloaded file using PGP signatures or hashes. The page is divided into sections for HTTP and BACKUP SITES, each with a link to the same download URL. A note at the bottom states: "The full listing of mirror sites is also available."

### Step 3:

Extract the Cassandra distribution into folder 1,2 and 3:



#### Step 4:

Now edit the code in each section in the file `conf/cassandra.yaml` and configure each of the Cassandra's nodes JMX port to point to different port numbers.

Open up `cassandra.bat` file for each of your nodes under the `bin` directory and look for the value:

#### 1<sup>st</sup> file:

```
INFO 19:28:24 OutboundTcpConnection using coalescing strategy DISABLED
INFO 19:28:32 No gossip backlog; proceeding
INFO 19:28:32 Netty using Java NIO event loop
INFO 19:28:32 Using Netty Version: [netty-buffer=netty-buffer-4.0.23.Final.2081
netty-codec-http=4.0.23.Final.208198c, netty-codec-socks=netty-codec-socks-4.0.23
andler=netty-handler-4.0.23.Final.208198c, netty-transport=netty-transport-4.0.
l.208198c, netty-transport-sctp=netty-transport-sctp-4.0.23.Final.208198c, net
INFO 19:28:32 Starting listening for CQL clients on /127.0.0.2:9042...
INFO 19:28:32 Binding thrift service to /127.0.0.2:9160
INFO 19:28:32 Listening for thrift clients...
```

#### 2<sup>nd</sup> file:

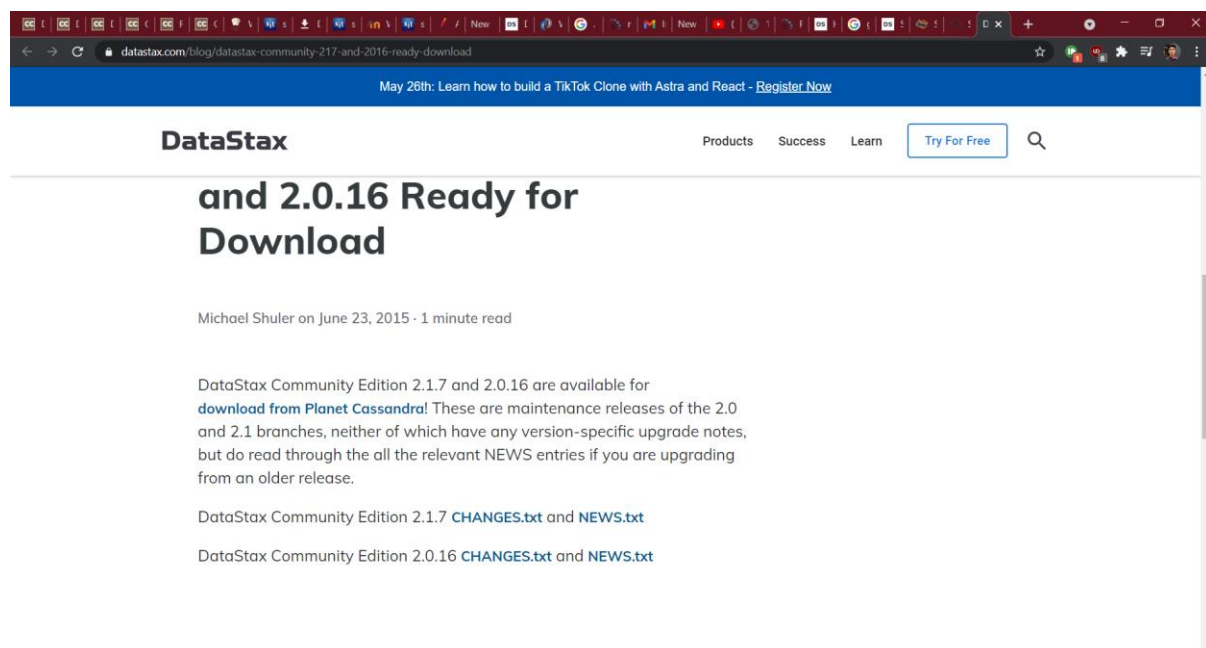
```
l.208198c, netty-transport-sctp=netty-transport-sctp-4.0.23.Final.208198c, netty-transport-u
INFO 19:28:32 Starting listening for CQL clients on /127.0.0.2:9042...
INFO 19:28:32 Binding thrift service to /127.0.0.2:9160
INFO 19:28:32 Listening for thrift clients...
INFO 19:34:00 Handshaking version with /127.0.0.3
INFO 19:34:00 Node /127.0.0.3 has restarted, now UP
INFO 19:34:00 Node /127.0.0.3 state jump to normal
INFO 19:34:00 InetAddress /127.0.0.3 is now UP
WARN 19:34:01 Not marking nodes down due to local pause of 339310798694 > 5000000000
INFO 19:34:22 Handshaking version with /127.0.0.4
INFO 19:34:24 Node /127.0.0.4 is now part of the cluster
INFO 19:34:24 InetAddress /127.0.0.4 is now UP
```

3<sup>rd</sup> file:

```
C:\opt\cassandra\1\bin>nodetool.bat status
Starting NodeTool
Datacenter: datacenter1
=====
Status=Up/Down
-/ State=Normal/Leaving/Joining/Moving
-- Address      Load       Tokens     Owns    Effective    Host ID                      Rack
UN  127.0.0.2    269.34 KB  256        69.3%           14df64f7-82f2-4bf2-a4ec-add12d60d121 rack1
UN  127.0.0.3    233.42 KB  256        62.8%           edd17770-58a9-4b2f-a8f7-8673d045a6a8 rack1
UN  127.0.0.4    107.59 KB  256        67.9%           1225ce9d-4de7-445c-9451-b415052045a3 rack1

C:\opt\cassandra\1\bin>
```

2. Install the DataStax OpsCenter community edition (<https://www.datastax.com/blog/datastax-community-217-and-2016-ready-download>) and configure it for above cluster formed.



3. Use Case - Weather Station IoT Temperature Sensor Data :  
There are set of weather stations at different remote location with “weatherStationID”. Each station record the

## 2018BCGRP24

temperature after every 5 minutes and push the data to nearest node in above cluster. Design the cluster database to hold these weather data. User should be able to retrieve the data in any dimensions

1) Creating table sensor\_data

```
Terminal +
Your Interactive Bash Terminal.
$
$ cqlsh
Connected to Cassandra Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 4.0-beta2 | CQL spec 3.4.5 | Native protocol v4]
Use HELP for help.
cqlsh>
cqlsh>
cqlsh> CREATE KEYSPACE sensor_data
... WITH replication = {
...   'class': 'NetworkTopologyStrategy',
...   'DC-Houston': 1 };
cqlsh> USE sensor_data;
cqlsh:sensor_data> CREATE TABLE networks (
...   bucket TEXT,
...   name TEXT,
...   description TEXT,
...   region TEXT,
...   num_sensors INT,
...   PRIMARY KEY ((bucket),name)
... );
cqlsh:sensor_data> CREATE TABLE temperatures_by_network (
...   network TEXT,
...   week DATE,
...   date_hour TIMESTAMP,
...   sensor TEXT,
...   avg_temperature FLOAT,
...   latitude DECIMAL,
...   longitude DECIMAL,
...   PRIMARY KEY ((network,week),date_hour,sensor)
... ) WITH CLUSTERING ORDER BY (date_hour DESC, sensor ASC);
```

```
cqlsh:sensor_data> CREATE TABLE sensors_by_network (
...   network TEXT,
...   sensor TEXT,
...   latitude DECIMAL,
...   longitude DECIMAL,
...   characteristics MAP<TEXT,TEXT>,
...   PRIMARY KEY ((network),sensor)
... );
cqlsh:sensor_data> CREATE TABLE temperatures_by_sensor (
...   sensor TEXT,
...   date DATE,
...   timestamp TIMESTAMP,
...   value FLOAT,
...   PRIMARY KEY ((sensor,date),timestamp)
... ) WITH CLUSTERING ORDER BY (timestamp DESC);
cqlsh:sensor_data> SOURCE '~/sensor_data.cql'
cqlsh:sensor_data> SELECT * FROM networks;
```

bucket	name	description	num_sensors	region
all	forest-net	forest fire detection network	3	south
all	volcano-net	volcano monitoring network	2	north

## 2018BCGRP24

### 2) Getting temperature of different zones.

```
cqlsh:sensor_data> SELECT * FROM temperatures_by_sensor;
```

sensor	date	timestamp	value
s1001	2020-07-04	2020-07-04 12:59:59.000000+0000	98
s1001	2020-07-04	2020-07-04 12:00:01.000000+0000	97
s1001	2020-07-04	2020-07-04 00:59:59.000000+0000	79
s1001	2020-07-04	2020-07-04 00:00:01.000000+0000	80
s1001	2020-07-05	2020-07-05 12:59:59.000000+0000	99
s1001	2020-07-05	2020-07-05 12:00:01.000000+0000	98
s1001	2020-07-05	2020-07-05 00:59:59.000000+0000	80
s1001	2020-07-05	2020-07-05 00:00:01.000000+0000	81
s1002	2020-07-06	2020-07-06 12:59:59.000000+0000	110
s1002	2020-07-06	2020-07-06 12:00:01.000000+0000	108
s1002	2020-07-06	2020-07-06 00:59:59.000000+0000	90
s1002	2020-07-06	2020-07-06 00:00:01.000000+0000	90
s1003	2020-07-04	2020-07-04 12:59:59.000000+0000	98
s1003	2020-07-04	2020-07-04 12:00:01.000000+0000	99
s1003	2020-07-04	2020-07-04 00:59:59.000000+0000	80
s1003	2020-07-04	2020-07-04 00:00:01.000000+0000	81
s1003	2020-07-06	2020-07-06 12:59:59.000000+0000	1429
s1003	2020-07-06	2020-07-06 12:00:01.000000+0000	1315
s1003	2020-07-06	2020-07-06 00:59:59.000000+0000	90
s1003	2020-07-06	2020-07-06 00:00:01.000000+0000	90
s1003	2020-07-05	2020-07-05 12:59:59.000000+0000	102
s1003	2020-07-05	2020-07-05 12:00:01.000000+0000	101

```
(36 rows)
cqlsh:sensor_data> SELECT name, description,
...           region, num_sensors
... FROM networks
... WHERE bucket = 'all';
```

name	description	region	num_sensors
forest-net	forest fire detection network	south	3
volcano-net	volcano monitoring network	north	2

```
(2 rows)
cqlsh:sensor_data>
```

### 3) Getting average temperature at specific time

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-06 12:00:00.000000+0000	106.5	30.526503	-95.582815	s1001
2020-07-06 12:00:00.000000+0000	109	30.518650	-95.583585	s1002
2020-07-06 12:00:00.000000+0000	1372	30.515056	-95.556225	s1003
2020-07-06 00:00:00.000000+0000	90.5	30.526503	-95.582815	s1001
2020-07-06 00:00:00.000000+0000	90	30.518650	-95.583585	s1002
2020-07-06 00:00:00.000000+0000	90.5	30.515056	-95.556225	s1003
2020-07-05 12:00:00.000000+0000	98.5	30.526503	-95.582815	s1001
2020-07-05 12:00:00.000000+0000	99.5	30.518650	-95.583585	s1002
2020-07-05 12:00:00.000000+0000	101.5	30.515056	-95.556225	s1003
2020-07-05 00:00:00.000000+0000	80.5	30.526503	-95.582815	s1001
2020-07-05 00:00:00.000000+0000	82	30.518650	-95.583585	s1002
2020-07-05 00:00:00.000000+0000	82.5	30.515056	-95.556225	s1003



## 2018BCGRP24

```
cqlsh:sensor_data> SELECT date_hour, avg_temperature,
... latitude, longitude, sensor
... FROM temperatures_by_network
... WHERE network = 'forest-net'
... AND week = '2020-06-28'
... AND date_hour >= '2020-07-04'
... AND date_hour < '2020-07-07';
```

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-04 12:00:00.000000+0000	97.5	30.526503	-95.582815	s1001
2020-07-04 12:00:00.000000+0000	100	30.518650	-95.583585	s1002
2020-07-04 12:00:00.000000+0000	98.5	30.515056	-95.556225	s1003
2020-07-04 00:00:00.000000+0000	79.5	30.526503	-95.582815	s1001
2020-07-04 00:00:00.000000+0000	81	30.518650	-95.583585	s1002
2020-07-04 00:00:00.000000+0000	80.5	30.515056	-95.556225	s1003

```
cqlsh:sensor_data> SELECT date_hour, avg_temperature,
... latitude, longitude, sensor
... FROM temperatures_by_network
... WHERE network = 'forest-net'
... AND week IN ('2020-07-05', '2020-06-28')
... AND date_hour >= '2020-07-04'
... AND date_hour < '2020-07-07';
```

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-04 12:00:00.000000+0000	97.5	30.526503	-95.582815	s1001
2020-07-04 12:00:00.000000+0000	100	30.518650	-95.583585	s1002
2020-07-04 12:00:00.000000+0000	98.5	30.515056	-95.556225	s1003
2020-07-04 00:00:00.000000+0000	79.5	30.526503	-95.582815	s1001
2020-07-04 00:00:00.000000+0000	81	30.518650	-95.583585	s1002
2020-07-04 00:00:00.000000+0000	80.5	30.515056	-95.556225	s1003
2020-07-06 12:00:00.000000+0000	106.5	30.526503	-95.582815	s1001
2020-07-06 12:00:00.000000+0000	109	30.518650	-95.583585	s1002
2020-07-06 12:00:00.000000+0000	1372	30.515056	-95.556225	s1003
2020-07-06 00:00:00.000000+0000	90.5	30.526503	-95.582815	s1001
2020-07-06 00:00:00.000000+0000	90	30.518650	-95.583585	s1002
2020-07-06 00:00:00.000000+0000	90.5	30.515056	-95.556225	s1003
2020-07-05 12:00:00.000000+0000	98.5	30.526503	-95.582815	s1001
2020-07-05 12:00:00.000000+0000	99.5	30.518650	-95.583585	s1002
2020-07-05 12:00:00.000000+0000	101.5	30.515056	-95.556225	s1003
2020-07-05 00:00:00.000000+0000	80.5	30.526503	-95.582815	s1001
2020-07-05 00:00:00.000000+0000	82	30.518650	-95.583585	s1002
2020-07-05 00:00:00.000000+0000	82.5	30.515056	-95.556225	s1003

## Conclusion:

Using DataStax we fetch the data and store it in Cassandra DB using clusters.

## References:

- <https://extendit.us/articles/steps-configure-multiple-nodes-cassandra-single-windows-machine>
- <https://www.datastax.com/blog/datastax-community-217-and-2016-ready-download>.